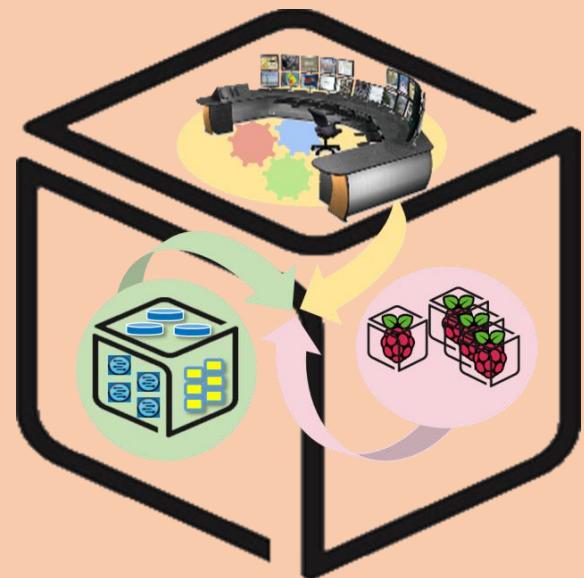


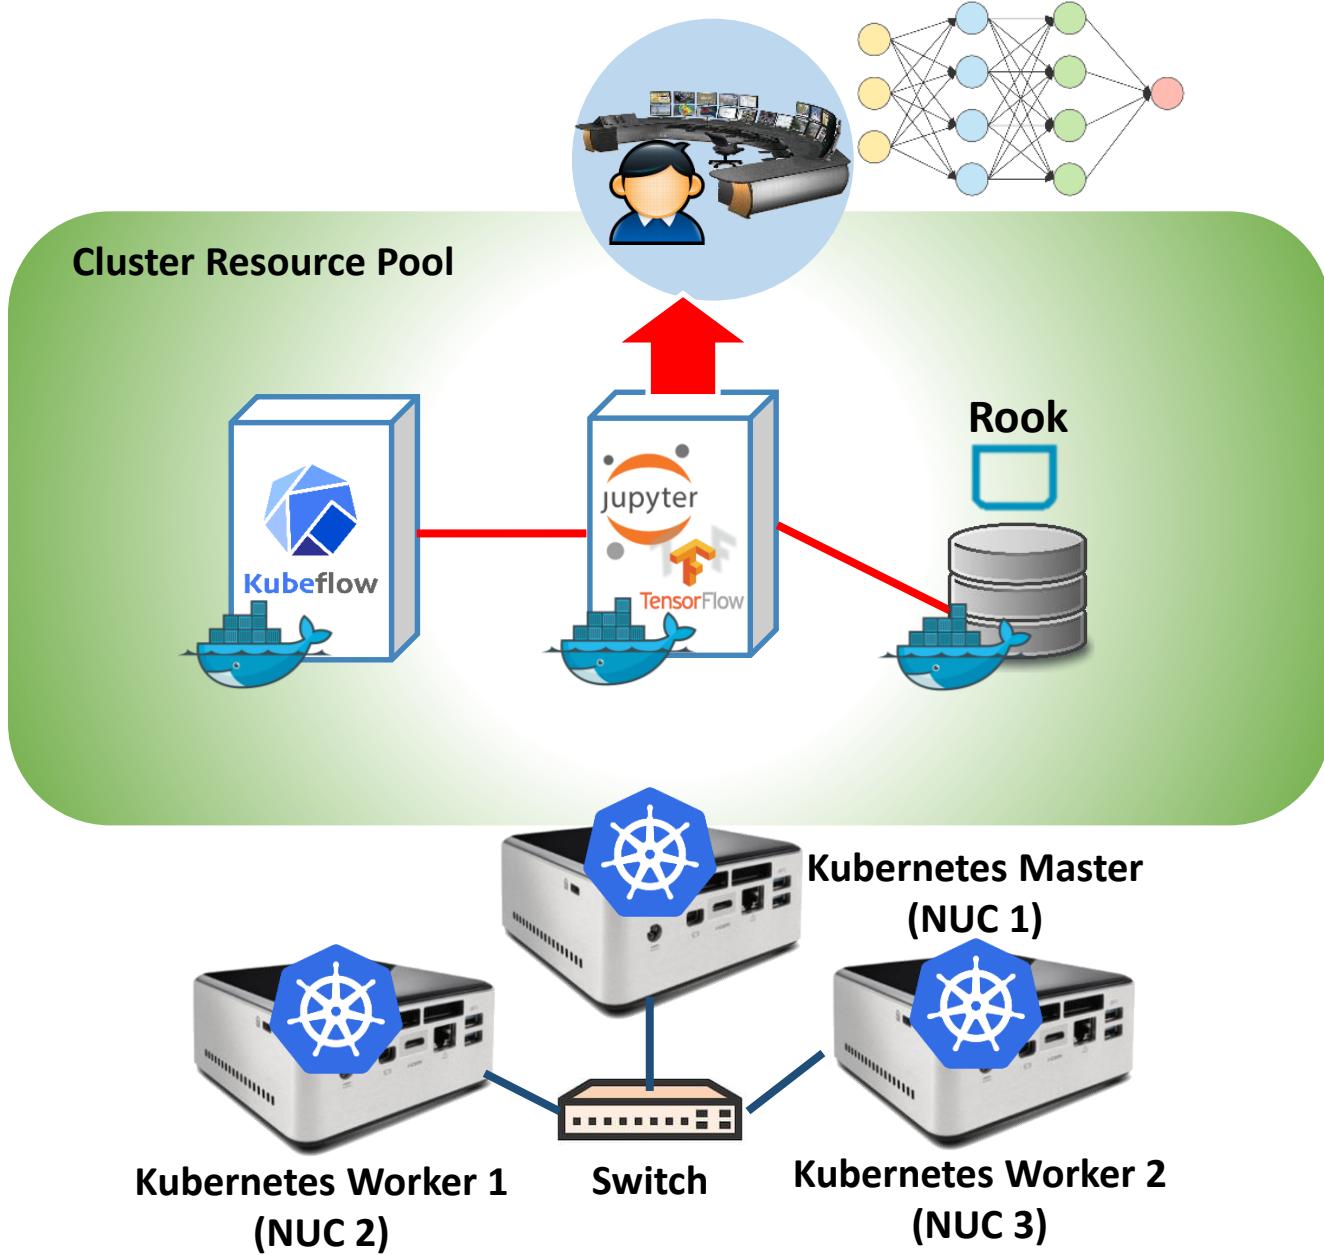
# Computer Systems For AI-inspired Cloud Theory & Lab.

## Lab #6: Analytics

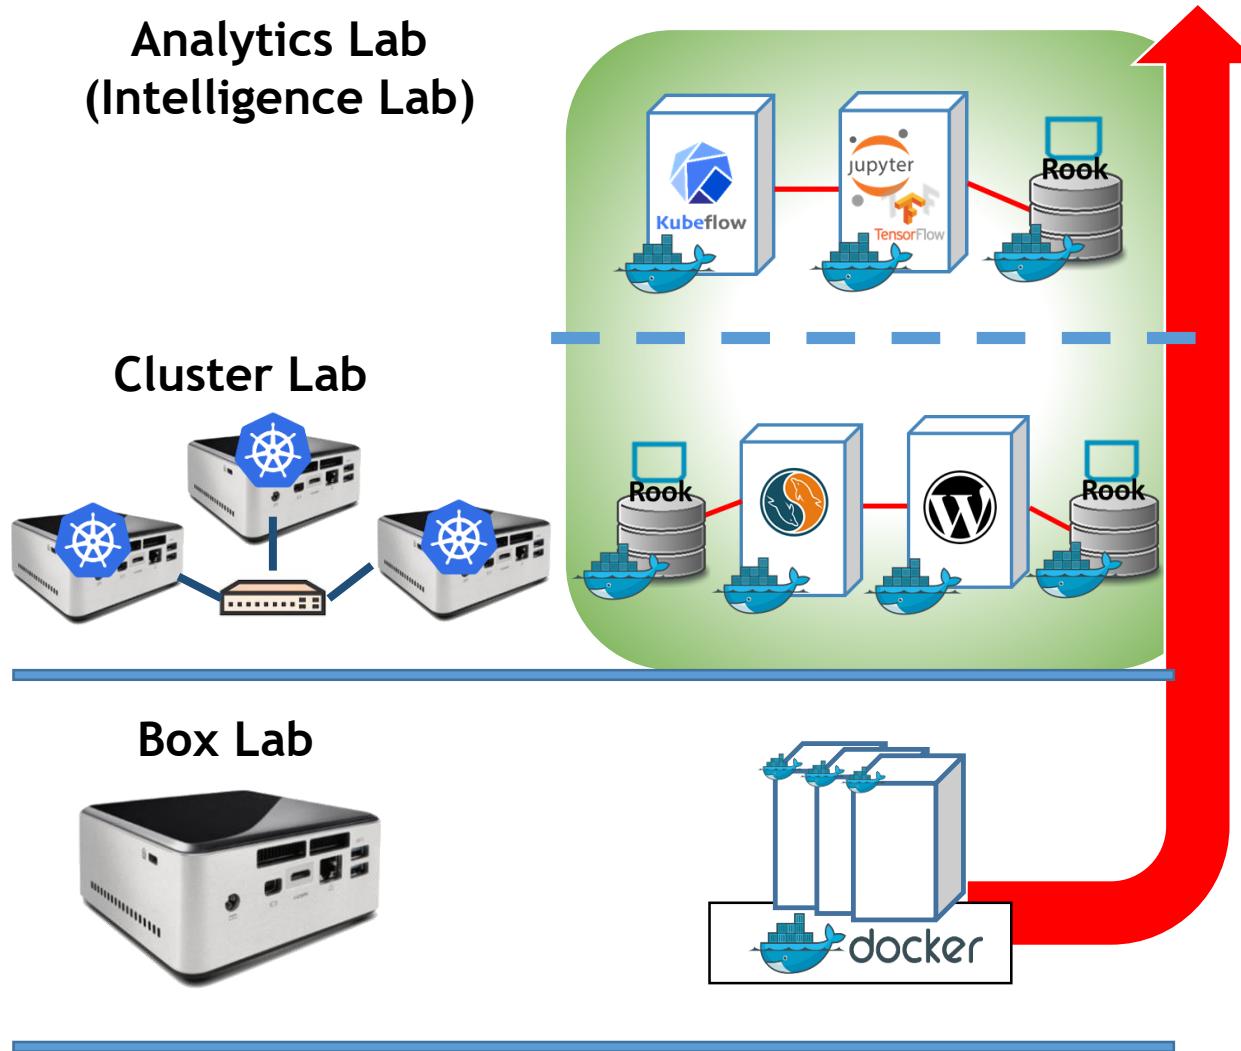


<https://github.com/SmartX-Labs/SmartX-Mini-MOOC>

# Analytics Lab: Concept



# SmartX Labs #1/#5/#6: Relationship

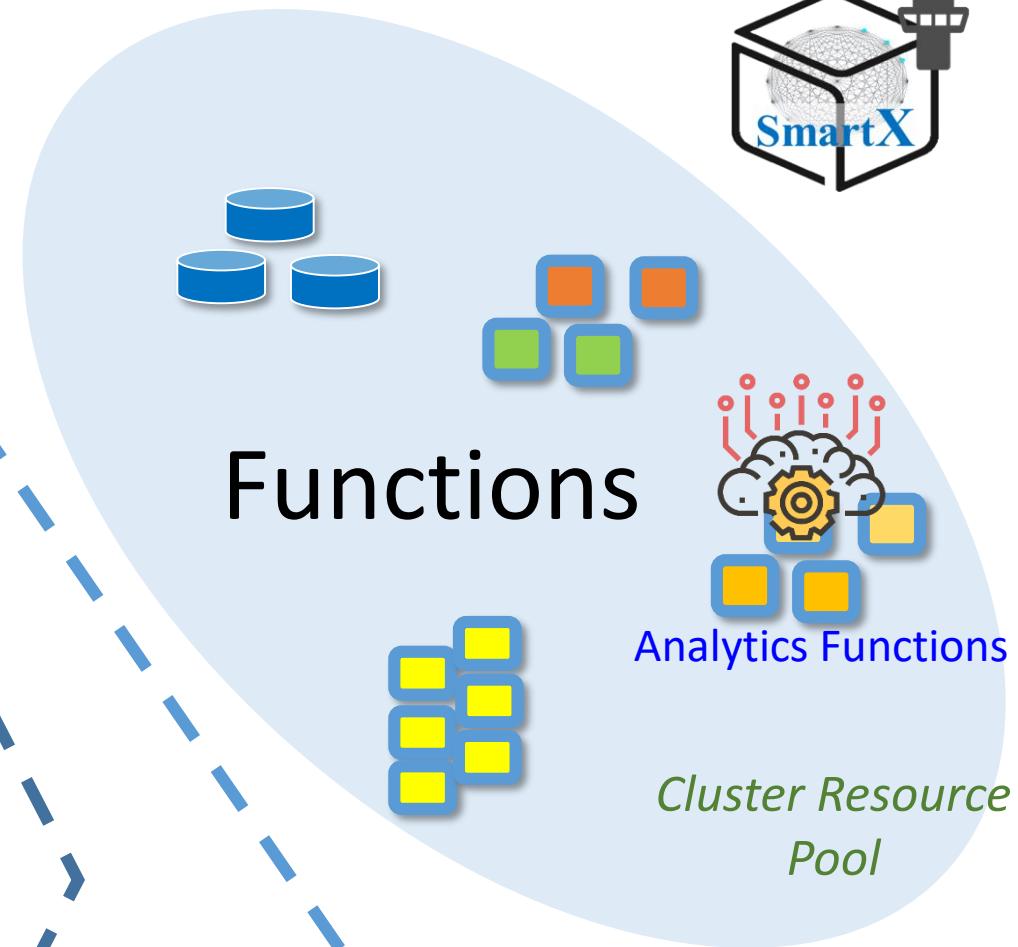
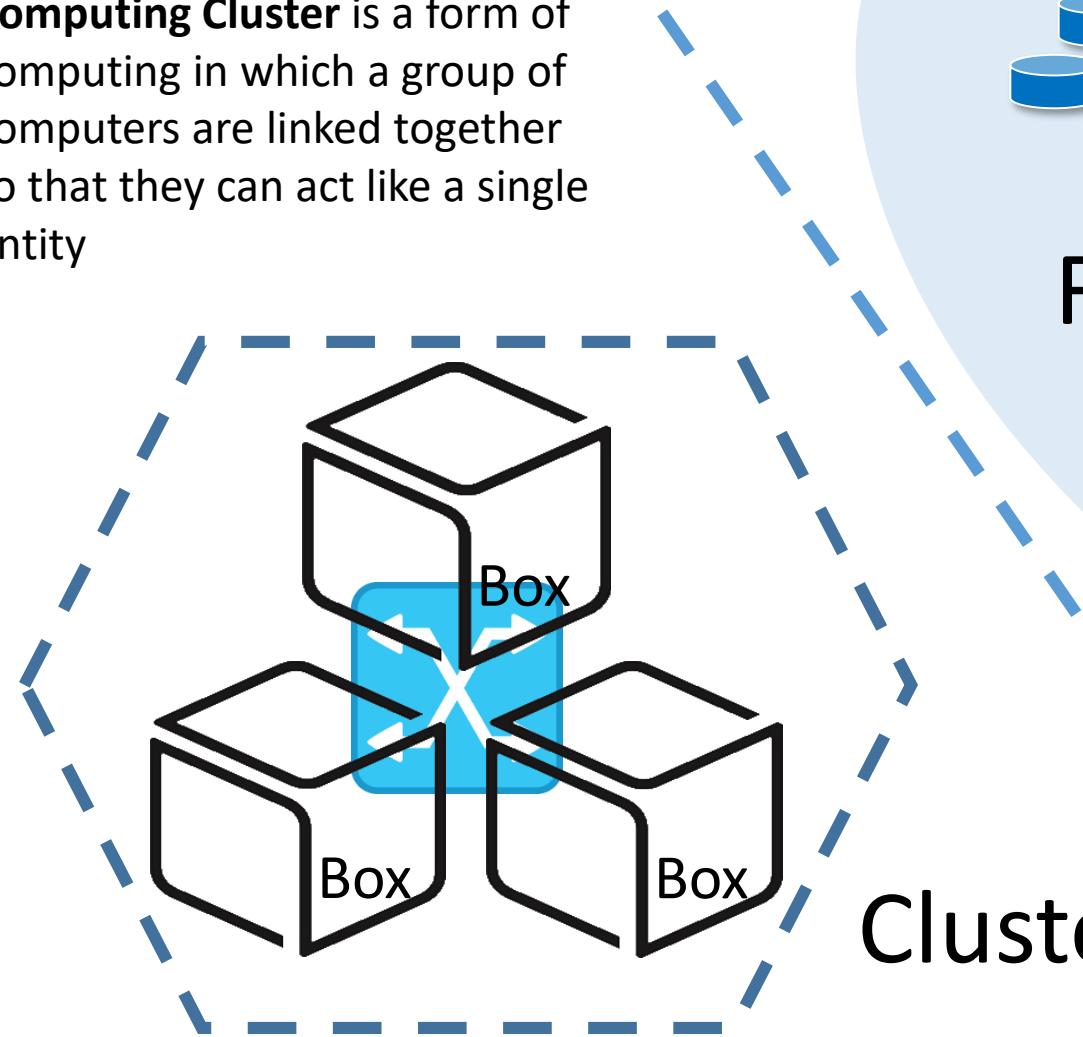


# Theory



# SmartX Cluster: Inter-connected SmartX Boxes

**Computing Cluster** is a form of computing in which a group of computers are linked together so that they can act like a single entity

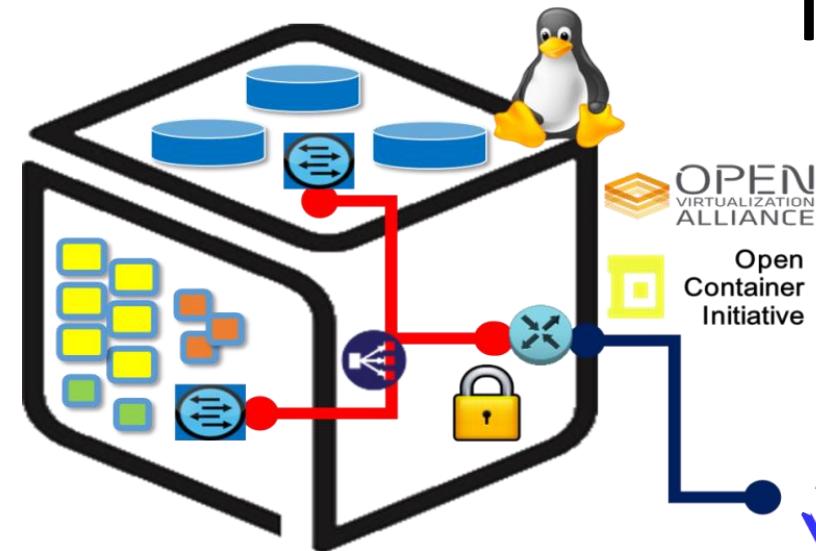




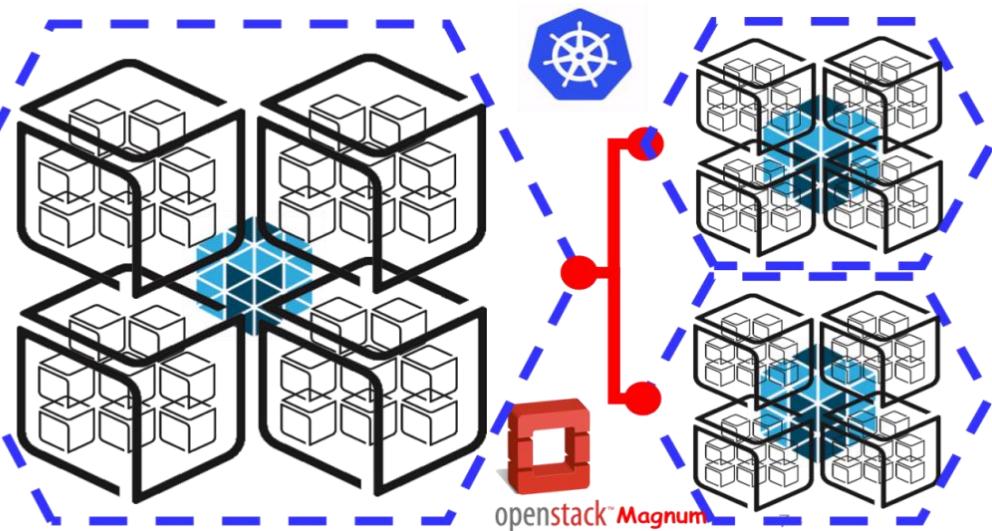
# Inter-Connected Functions inside a Box & across Clusters

**p+v+c Harmonization Challenge:**  
**p(Baremetal) + v(VM) + c(Container)**

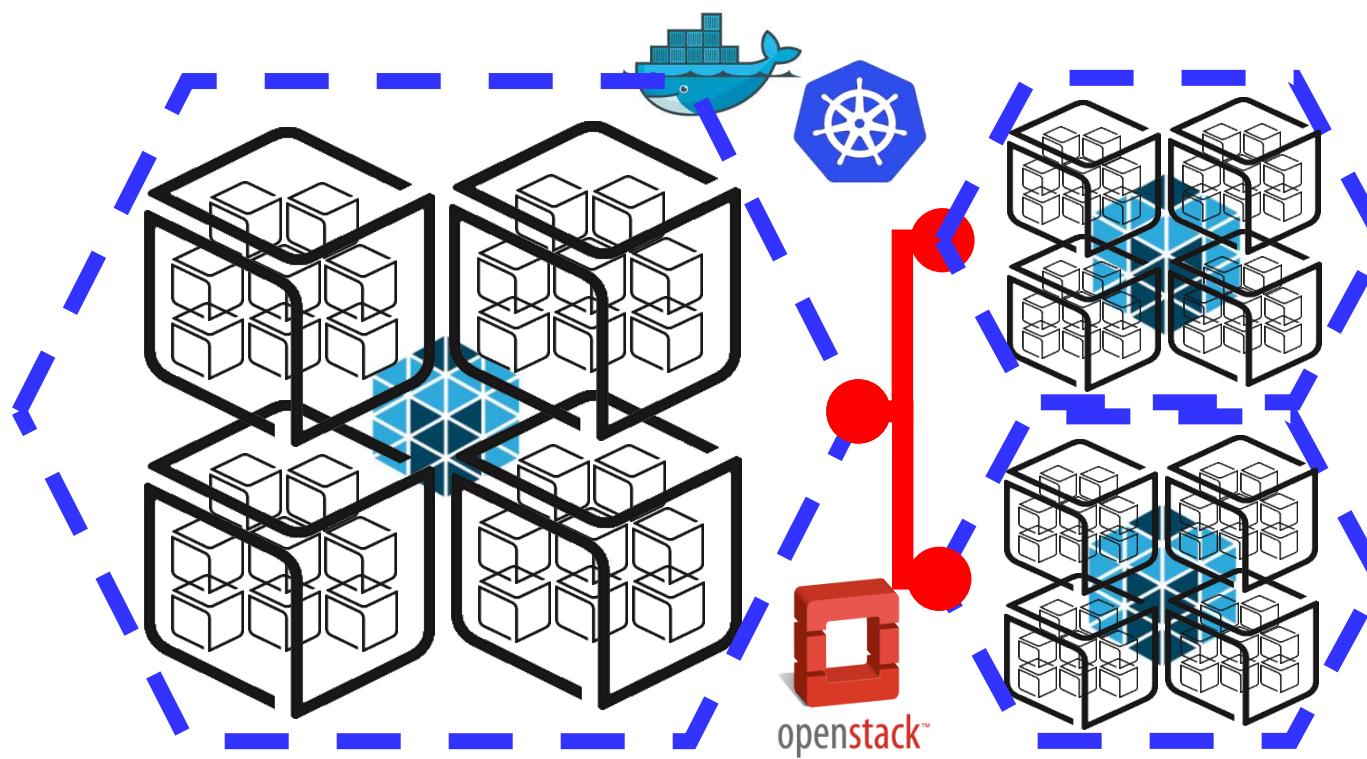
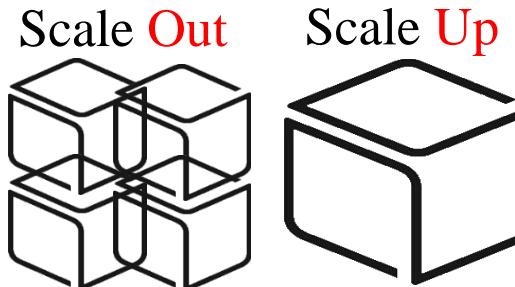
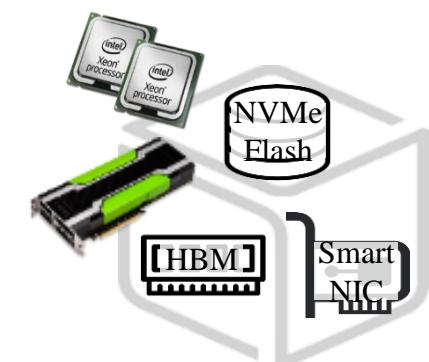
Inside a Box



Across Clusters

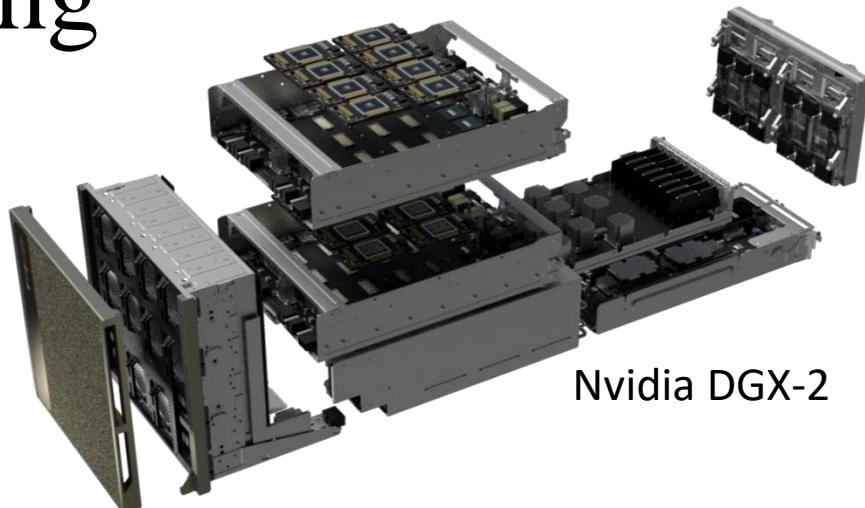


# Computer System: Resource Scaling/Pooling with Clustering

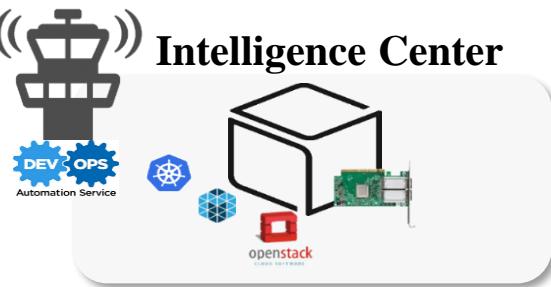


# Cluster for AI Computing

HPC + HPDA (BigData)  
 → AI (ML/DL)



## Intelligence Center



## Multi-node AI Computing Cluster with Optimized DL Tools

### GPU Acceleration

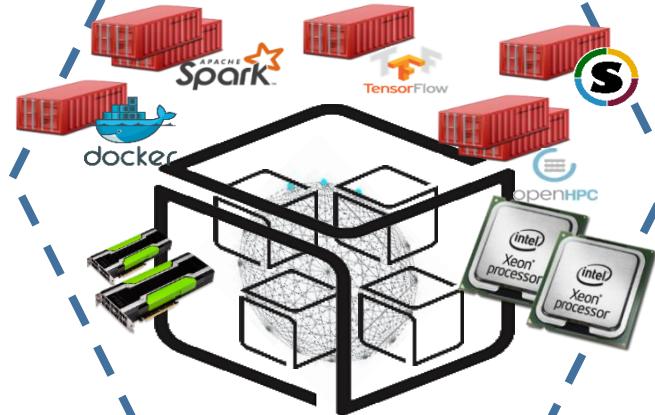


### Parallel File System

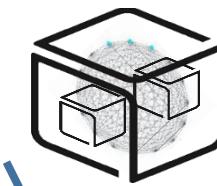


### Container Orchestration

kubernetes



ceph



### Cloud Storage Cluster

### Cloud Storage

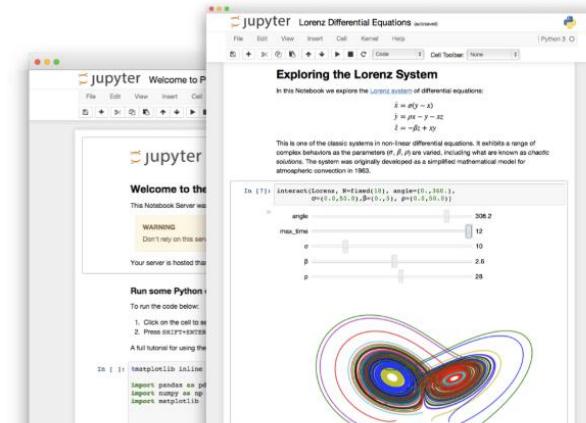


# Machine Learning: TensorFlow & Jupyter Notebook



**TensorFlow** is an **open-source machine learning library** for research and production. TensorFlow offers APIs for beginners and experts to develop for desktop, mobile, web, and cloud.

<https://github.com/tensorflow/tensorflow>

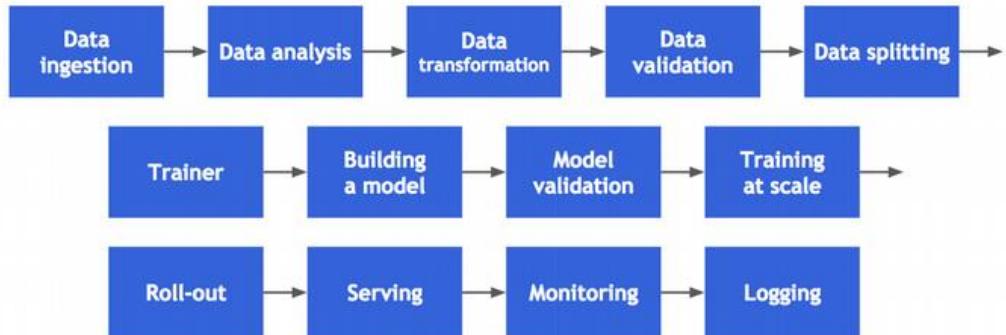


**Jupyter Notebook** is an open-source web application that allows you to create and share documents that contain live code, equations, visualizations and narrative text. Uses include: data cleaning and transformation, numerical simulation, statistical modeling, data visualization, **machine learning**, and much more.

# Machine Learning: Kubeflow



**Kubeflow**



The **Kubeflow** project is dedicated to making deployments of machine learning (ML) workflows on Kubernetes simple, portable and scalable. Our goal is not to recreate other services, but to provide a straightforward way to deploy best-of-breed open-source systems for ML to diverse infrastructures. Anywhere you are running Kubernetes, you should be able to run Kubeflow.



Notebooks



TesorFlow  
model Training



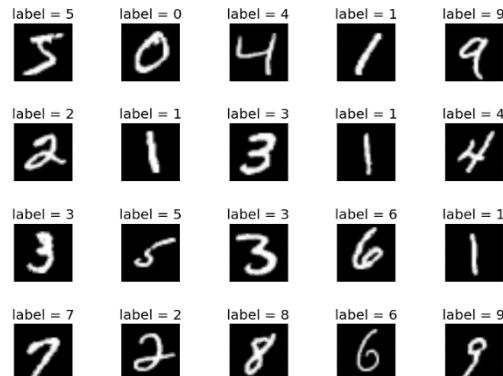
Model serving



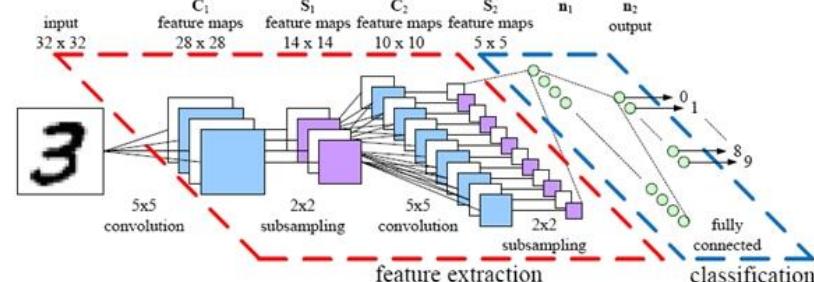
Multi-ML framework

# MNIST handwritten digit Classification

0 0 0 0 0 0 0 0 0 0 0 0 0 0 0  
 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1  
 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2  
 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3  
 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4  
 5 5 5 5 5 5 5 5 5 5 5 5 5 5 5  
 6 6 6 6 6 6 6 6 6 6 6 6 6 6 6  
 7 7 7 7 7 7 7 7 7 7 7 7 7 7 7  
 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8  
 9 9 9 9 9 9 9 9 9 9 9 9 9 9 9



The MNIST database (Modified National Institute of Standards and Technology database) is a large database of handwritten digits that is commonly used for training various image processing systems.[1][2] The database is also widely used for training and testing in the field of machine learning.

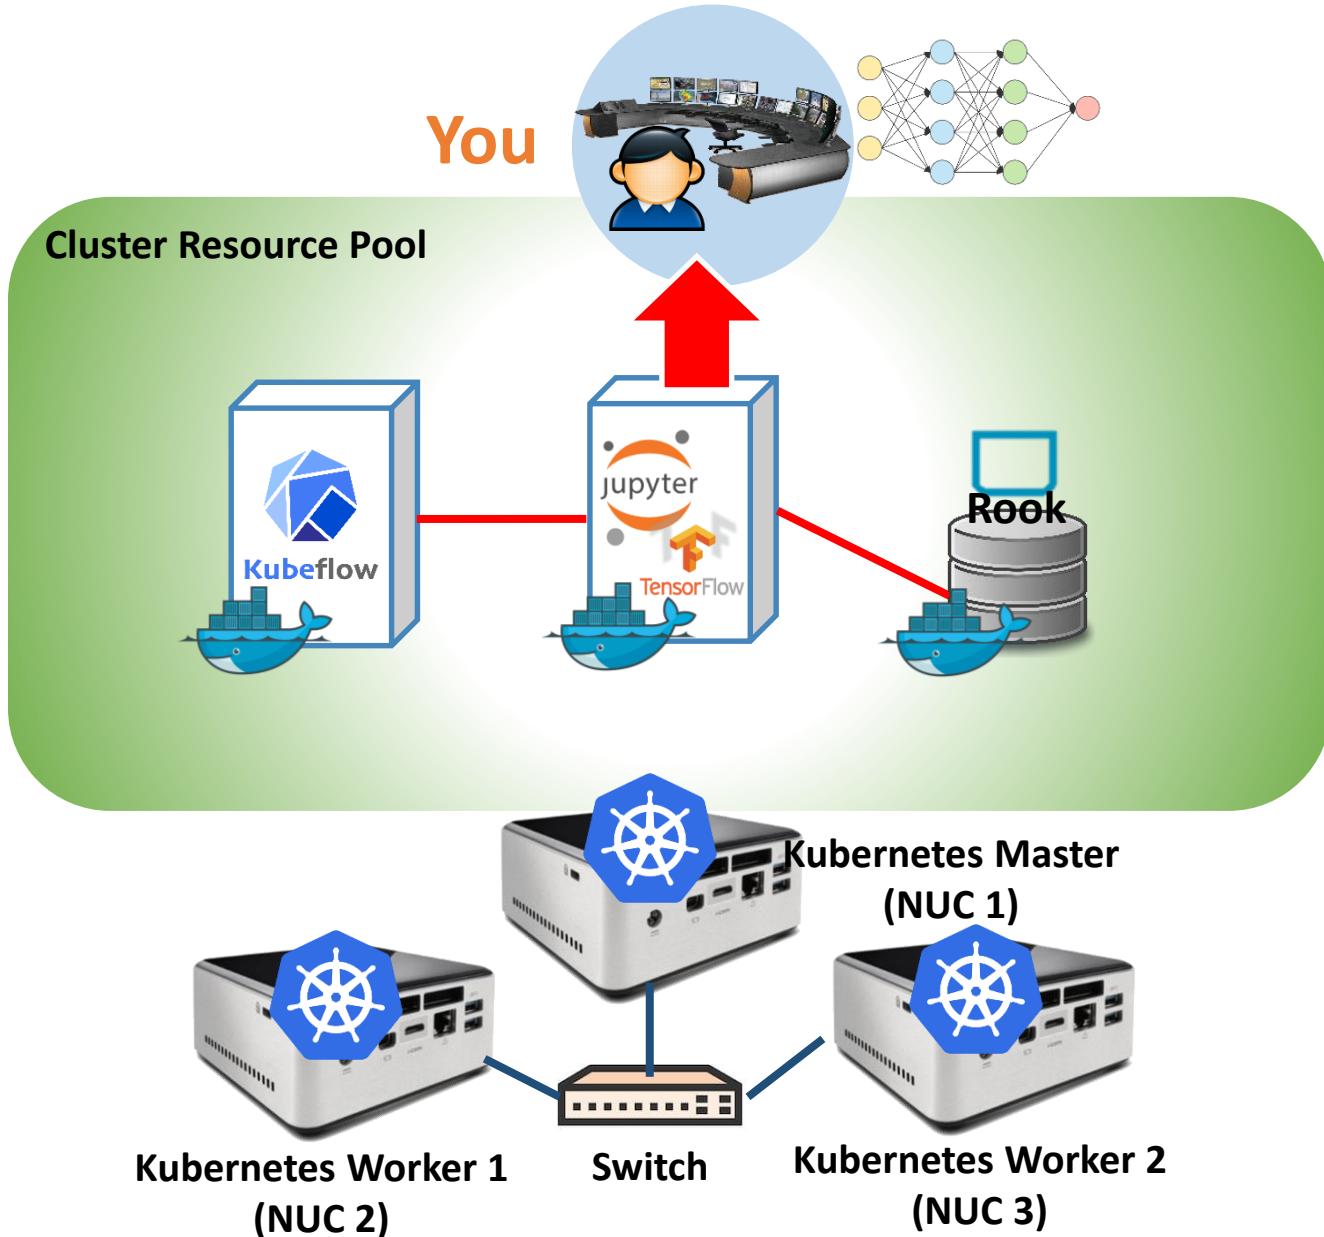


# Practice



# Analytics Lab: Concept

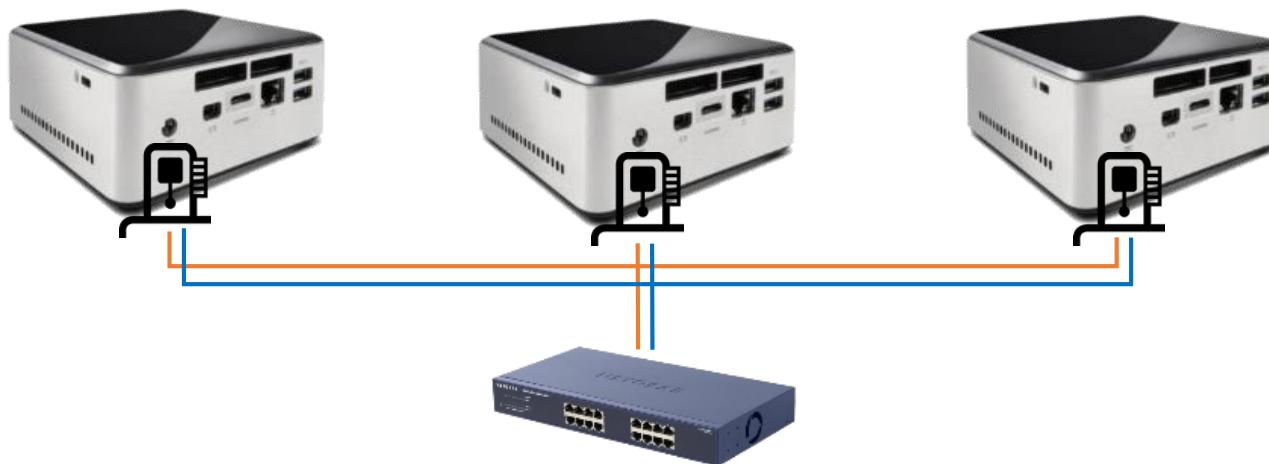
Lab #6: Analytics 13



# #0 - Lab Preparation (1/2)

## Wired connection

**NAME:** NUC5i5MYHE (NUC PC)  
**CPU:** i5-5300U @2.30GHz  
**CORE:** 4  
**Memory:** 16GB DDR3  
**HDD:** 94GB



**NAME:** netgear prosafe 16 port gigabit switch(Switch)  
**Network Ports:** 16 auto-sensing 10/100/1000 Mbps Ethernet ports

# #0 - Lab Preparation (2/2)

- Check your cluster is running healthy

For **NUC1**

```
$ kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
nuc01	Ready	master	10d	v1.11.2
nuc02	Ready	<none>	10d	v1.11.2
nuc03	Ready	<none>	10d	v1.11.2



Check all nodes are ready.

```
$ kubectl get pods -n rook-ceph
```

NAME	READY	STATUS	RESTARTS
rook-ceph-mgr-a-9c44495df-lfs4m	1/1	Running	0
rook-ceph-mon0-5j655	1/1	Running	0
rook-ceph-mon1-rkggs	1/1	Running	0
rook-ceph-mon2-vvp2n	1/1	Running	0
rook-ceph-osd-id-0-8694878c4b-9zz5l	1/1	Running	0
rook-ceph-osd-id-1-756995f97b-9hdhk	1/1	Running	0
rook-ceph-osd-prepare-nuc02-26nj6	0/1	Completed	0
rook-ceph-osd-prepare-nuc03-cf49p	0/1	Completed	0

Check Rook are running healthy on your cluster.

# #1-1 Cluster Preparations for ML: Kubeflow Installation (1/4)

For **NUC1**



- Set Rook Storageclass to default for kubeflow

```
$ kubectl patch storageclass rook-ceph-block -p \  
> '{"metadata": {"annotations":{"storageclass.kubernetes.io/is-default-class":"true"}}}'
```

- Check

```
$ kubectl get storageclasses
```

# #1-1 Cluster Preparations for ML: Kubeflow Installation (2/4)



For **NUC1**

## - Install Kubeflow

### //configuration

```
$ export BASE_DIR=$HOME/kubeflow  
$ mkdir -p $BASE_DIR && pushd $BASE_DIR
```

### // Download kfctl

```
$ KFCTL_FILE=kfctl_v0.7.1-2-g55f9b2a_linux.tar.gz  
$ wget https://github.com/kubeflow/kfctl/releases/download/v0.7.1/${KFCTL_FILE}  
$ tar -xf $KFCTL_FILE && rm $KFCTL_FILE  
$ sudo mv kfctl /usr/local/bin/kfctl
```

Continued on next page



# #1-1 Cluster Preparations for ML: Kubeflow Installation (3/4)

For **NUC1**

## - Install Kubeflow (cont'd)

*// Set environment variables*

```
$ export KF_NAME="my-kubeflow"  
$ export KF_DIR=${BASE_DIR}/${KF_NAME}  
$ export PATH=$PATH:$BASE_DIR  
$ export CONFIG_URI="https://raw.githubusercontent.com/kubeflow\  
/manifests/v0.7- branch/kfdef/kfctl_k8s_istio.0.7.0.yaml"
```

*// Install kubeflow*

```
$ rm -rf $KF_DIR; mkdir -p $KF_DIR && cd $KF_DIR  
$ kfctl apply -V -f $CONFIG_URI  
$ popd
```



# #1-1 Cluster Preparations for ML: Kubeflow Installation (4/4)

For **NUC1**

- Check Kubeflow is running healthy**

```
$ watch kubectl get pods -n kubeflow
```

```
$ watch kubectl get svc -n istio-system
```

- Check the exposed port to access Jupyter hub**

```
$ kubectl get services -n kubeflow
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
ambassador	ClusterIP	10.111.165.80	<none>	80/TCP
ambassador-admin	ClusterIP	10.101.217.43	<none>	8877/TCP
argo-ui	NodePort	10.101.8.20	<none>	80:30681/TCP
centraldashboard	ClusterIP	10.105.124.82	<none>	80/TCP
k8s-dashboard	ClusterIP	10.110.111.206	<none>	443/TCP
modeldb-backend	ClusterIP	10.104.50.63	<none>	6543/TCP
modeldb-db	ClusterIP	10.103.246.44	<none>	27017/TCP
modeldb-frontend	ClusterIP	10.102.138.220	<none>	3000/TCP
statsd-sink	ClusterIP	10.106.80.124	<none>	9002/TCP
tf-hub-0	ClusterIP	None	<none>	8000/TCP
tf-hub-lb	NodePort	10.107.131.150	<none>	80:32290/TCP
tf-job-dashboard	ClusterIP	10.111.220.254	<none>	80/TCP
vizier-core	NodePort	10.96.183.97	<none>	6789:30678/TCP
vizier-db	ClusterIP	10.99.58.20	<none>	3306/TCP
vizier-suggestion-bayesianoptimization	ClusterIP	10.98.249.26	<none>	6789/TCP
vizier-suggestion-grid	ClusterIP	10.100.145.210	<none>	6789/TCP
vizier-suggestion-hyperband	ClusterIP	10.108.171.180	<none>	6789/TCP
vizier-suggestion-random	ClusterIP	10.97.121.224	<none>	6789/TCP

You can access Jupyter hub at this address

[http://nuc01\\_IP:Exposed\\_port](http://nuc01_IP:Exposed_port)

Port 확인:

```
kubectl -n istio-system get svc istio-ingressgateway -o jsonpath='{.spec.ports[?(@.name=="http2")].nodePort}'
```

# #1-2 Deploy a ML Container: Create a Jupyter Notebook (1/3)



For **NUC1**

Open a web browser and enter the Jupyter hub address

[http://nuc01\\_IP:Exposed\\_port](http://nuc01_IP:Exposed_port)

The screenshot shows a Mozilla Firefox browser window with the title "JupyterHub - Mozilla Firefox". The address bar displays the URL "203.237.53.71:32290/hub/login". The main content is a "Sign in" form for JupyterHub. A yellow warning message at the top reads: "Warning: JupyterHub seems to be served over an unsecured HTTP connection. We strongly recommend enabling HTTPS for JupyterHub." Below the message, there are two input fields: "Username:" containing "netcs" and "Password:". A red box highlights the "Sign In" button at the bottom of the form.

Enter your username and click 'Sign In' button (you don't have to enter a password)

# #1-2 Deploy a ML Container: Create a Jupyter Notebook (2/3)



For **NUC1**

In Spawner page, you can choose container image and the size of resources to create a Jupyter Notebook container for Machine Learning

Select or enter the options as below

**Image:** gcr.io/kubeflow-images-public/tensorflow-1.9.0-notebook-cpu:v0.3.1

**CPU:** 2

**Memory:** 2Gi

jupyter Home Token Logout

### Spawner Options

Image	gcr.io/kubeflow-images-public/tensorflow-1.9.0-notebook-cpu
Advanced	
CPU	2.0
Memory	2Gi
Extra Resource Limits	{"nvidia.com/gpu": 3}
<b>Spawn</b>	

jupyter Home Token Logout

Your server is starting up.  
You will be redirected automatically when it's ready for you.



Spawning server...  
Event log

Click Spawn Button and you need to wait for a while

# #1-2 Deploy a ML Container: Create a Jupyter Notebook (3/3)



For **NUC1**

You can see your Jupyter notebook container is deployed as a pod on cluster

\$ kubectl get pods –n kubeflow

NAME	READY	STATUS	RESTARTS
ambassador-c97f7b448-998gd	3/3	Running	0
ambassador-c97f7b448-c2j2b	3/3	Running	0
ambassador-c97f7b448-dvlkv	3/3	Running	0
argo-ui-7495b79b59-l5xh6	1/1	Running	0
centraldashboard-798f8d68d5-5wshd	1/1	Running	0
jupyter-netscs	1/1	Running	0
jupyter-test	1/1	Running	0
modeldb-backend-d69695b66-99cqm	1/1	Running	0
modeldb-db-975db58f7-prpzf	1/1	Running	0

Now your Jupyter notebook is created...

The screenshot shows the Jupyter Notebook interface. At the top, there's a navigation bar with a logo, 'Logout', and 'Control Panel'. Below it, a menu bar has tabs for 'Files' (which is selected), 'Running', and 'Clusters'. A message says 'Select items to perform actions on them.' Below this is a file browser with a sidebar showing a folder structure: '0' and 'work'. The main area shows files with columns for 'Name', 'Last Modified', and 'File size'. One file is shown with a timestamp of '17시간 전'.

# #2-1 Running Analytics code: Running a Sample ML Code (1/3)



For **NUC01**

Download sample notebook including MNIST Machine Learning Code

\$ git clone <https://github.com/aymericdamien/TensorFlow-Examples/>

Or Download from web browser as below. (you need to unzip the file)

I and Examples for Beginners with Latest APIs <https://tensorflow.org>

examples deep-learning python machine-learning

its 4 branches 0 releases 50 contributors View license

New pull request Find file Clone or download ▾

TensorFlow-Examples-master\notebooks\3\_NeuralNetwork\convolutional\_network.ipynb

We will upload the sample notebook on your Jupyter and run it.  
The notebook include MNIST machine learning example code.

<https://github.com/aymericdamien/TensorFlow-Examples/>

# #2-1 Running Analytics code: Running a Sample ML Code (2/3)



For **NUC1**

jupyter

Logout Control Panel

Files Running Clusters

Select items to perform actions on them.

Upload New

0	/	Name <input type="button"/>	Last Modified <input type="button"/>	File size <input type="button"/>
17시간 전				
<input type="checkbox"/> work				

Press upload button to upload the sample notebook

jupyter

Logout Control Panel

Files Running Clusters

Select items to perform actions on them.

Upload New

0	/	Name <input type="button"/>	Last Modified <input type="button"/>	File size <input type="button"/>
18시간 전				
<input type="checkbox"/> work				
<input checked="" type="checkbox"/> convolutional_network.ipynb				몇 초 전 33.2 kB

Click it to open the notebook

Remember! we will use this file  
TensorFlow-Examples-  
master\notebooks\3\_NeuralNetwork\convolutio  
nal\_network.ipynb

# #2-1 Running Analytics code: Running a Sample ML Code (3/3)



For **NUC1**

The screenshot shows a Jupyter Notebook window titled 'jupyter convolutional\_network'. The 'Kernel' menu is open, and the 'Restart & Clear Output' option is highlighted with a red box. The notebook content displays a TensorFlow example for a convolutional network.

Now, you will run the MNIST example code in sample notebook.  
Click kernel → “Restart&Clear Output” button → “Restart&Run All” button

```
In [4]: # Define the input function for training
input_fn = tf.estimator.inputs.numpy_input_fn(
    x={'images': mnist.train.images}, y=mnist.train.labels,
    batch_size=batch_size, num_epochs=None, shuffle=True)
# Train the Model
model.train(input_fn, steps=num_steps)

INFO:tensorflow:Calling model_fn.
INFO:tensorflow:Done calling model_fn.
INFO:tensorflow>Create CheckpointSaverHook.
INFO:tensorflow:Graph was finalized.
INFO:tensorflow:Running local_init_op.
INFO:tensorflow:Done running local_init_op.
INFO:tensorflow:Saving checkpoints for 0 into /tmp/tmpo9DvpG/model.ckpt.
INFO:tensorflow:loss = 2.3231235, step = 1
INFO:tensorflow:global_step/sec: 7.88675
INFO:tensorflow:loss = 0.09013237, step = 101 (12.681 sec)
INFO:tensorflow:global_step/sec: 7.96045
INFO:tensorflow:loss = 0.087195996, step = 201 (12.562 sec)
INFO:tensorflow:global_step/sec: 7.93652
INFO:tensorflow:loss = 0.07184037, step = 301 (12.600 sec)
INFO:tensorflow:global_step/sec: 7.92234
INFO:tensorflow:loss = 0.15338697, step = 401 (12.622 sec)
```

The training takes a few minutes.

# #2-2 Running Analytics code: Check ML Training results



For **NUC1**

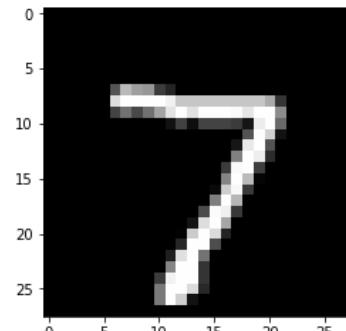
```
In [7]: # Evaluate the Model
# Define the input function for evaluating
input_fn = tf.estimator.inputs.numpy_input_fn(
    x={'images': mnist.test.images}, y=mnist.test.labels,
    batch_size=batch_size, shuffle=False)
# Use the Estimator 'evaluate' method
model.evaluate(input_fn)

INFO:tensorflow:Calling model_fn.
INFO:tensorflow:Done calling model_fn.
INFO:tensorflow:Starting evaluation at 2018-11-25-07:45:59
INFO:tensorflow:Graph was finalized.
INFO:tensorflow:Restoring parameters from /tmp/tmpp9DVsG/model.ckpt-2000
INFO:tensorflow:Running local_init_op.
INFO:tensorflow:Done running local_init_op.
INFO:tensorflow:Finished evaluation at 2018-11-25-07:46:03
INFO:tensorflow:Saving dict for global step 2000: accuracy = 0.9892, global_step = 2000, loss = 0.035650674
INFO:tensorflow:Saving 'checkpoint_path' summary for global step 2000: /tmp/tmpp9DVsG/model.ckpt-2000

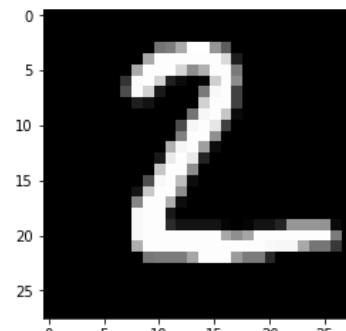
Out[7]: {'accuracy': 0.9892, 'global_step': 2000, 'loss': 0.035650674}
```

## Check training results

Your model has 98.92% accuracy!



Model prediction: 7



Model prediction: 2

Your Machine  
Learning model  
correctly identified  
the number in the  
images!

# Review

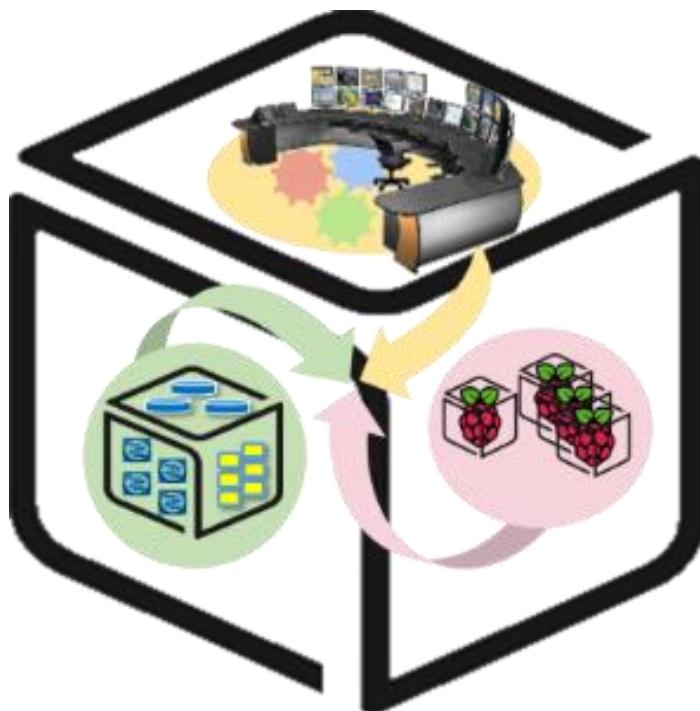


# Lab Summary

With Analytics (Intelligence) Lab, you have experimented

1. How to create ML/DL environment on a container-orchestrated cluster? (Kubeflow, ...)
2. How to operate desired ML training by testing selected ML code (i.e., neural networks) over the prepared training data?
3. Do you understand the overall workflow for running ML/DL?

# Thank You for Your Attention Any Questions?



mini@smartx.kr